

Human Emotion Recognition

Advanced Programming and Deep Learning for AI

Georgii Kutivadze



[Source](#)

Outline

1

Introduction



1. What is the goal?
2. What are the tasks?

2

Data & Preprocessing



1. What Datasets are used?
2. What is the Size & Image Format?
3. What is the Class Distribution?
4. Which Role each Dataset plays?

3

Models & Training



1. What is the Model Architecture
2. How many epochs? What is the Learning Rate?
3. What are the Train and Validation Loss?
4. What are the Train and Validation Accuracy?

4

Evaluation & Comparison



1. Which Model is the Best?
2. What is the Best Model's Performance on Test Set?
3. Is the Score Satisfactory?

5

Generalization Check



1. What if we apply zero-shot of our Best Model on a New Dataset?
2. What if we Fine-Tune our Best Model on it and Repeat Evaluation?
3. What is the Generalization Capability of our Best Model?

6

Conclusion



1. What are the key findings?
2. What are the most promising directions for future work?


Can We Recognize Human Emotions from a Face Image?

GOAL

Systematically compare baseline and advanced deep-learning models for **facial expression recognition** on FER-2013, and study how well the best model transfers to RAF-DB

SUBTASKS

1. Build and train **baseline models** on FER-2013
2. Design and train stronger **CNN architectures** on FER-2013
3. Compare all models on the same **FER-2013** train/val/test splits and select the **best one**
4. Evaluate **transfer** of the best FER model to **RAF-DB** (before and after fine-tuning)
5. Analyze class distributions and **errors** on both datasets to understand **remaining limitations**



[Source](#)

Datasets: FER-2013 for Main Training, RAF-DB for Fine-Tuning & Generalization


SIZE

- **FER-2013:** larger (~35k images)
- **RAF-DB:** smaller (~15k images)

1010
1010

IMAGE FORMAT

- **FER-2013:** low-resolution, grayscale, 48×48
- **RAF-DB:** higher-quality RGB images, resized and converted to grayscale in our pipeline




LABEL SPACE

- Both use the same **7 basic emotions:** Angry, Disgust, Fear, Happy, Sad, Surprise, Neutral
- For RAF-DB, class indices are **remapped** to match FER-2013 labels




CLASS DISTRIBUTION

- Both are **heavily skewed** towards Happy, Neutral, and Sad
- Minority classes (Disgust, Fear) are especially **under-represented** in RAF-DB



ROLE IN THIS PROJECT

- **FER-2013:** main dataset for training and model comparison
- **RAF-DB:** used for **transfer learning** and checking cross-dataset generalization





Model №1: Simple MLP

ARCHITECTURE


```
class SimpleMLP(nn.Module):
    def __init__(self, num_classes=7):
        super().__init__()
        self.flatten = nn.Flatten()
        self.net = nn.Sequential(
            nn.Linear(1 * 48 * 48, 256),
            nn.ReLU(),
            nn.Linear(256, num_classes)
        )

    def forward(self, x):
        x = self.flatten(x)
        x = self.net(x)
        return x
```



TRAIN LOSS

1.5884




TRAIN ACCURACY

38.24%



VAL LOSS

1.5726



VAL ACCURACY

39.54%



Model №2: Improved MLP (BatchNorm + Dropout)



ARCHITECTURE



```
class ImprovedMLP(nn.Module):
    def __init__(self, num_classes=7):
        super().__init__()
        self.flatten = nn.Flatten()
        self.net = nn.Sequential(
            nn.Linear(1 * 48 * 48, 512),
            nn.BatchNorm1d(512),
            nn.ReLU(),
            nn.Dropout(0.3),

            nn.Linear(512, 256),
            nn.BatchNorm1d(256),
            nn.ReLU(),
            nn.Dropout(0.3),

            nn.Linear(256, 128),
            nn.BatchNorm1d(128),
            nn.ReLU(),


            nn.Linear(128, num_classes)
        )

    def forward(self, x):
        x = self.flatten(x)
        x = self.net(x) # [B, 1, 48, 48] -> [B, 2304]# [B, 7]
        return x
```



TRAIN LOSS

1.4097




TRAIN ACCURACY

45.36%



VAL LOSS

1.4192



VAL ACCURACY

45.78%



Model №3: Simple CNN


ARCHITECTURE

```
class SimpleCNN(nn.Module):
    def __init__(self, num_classes=7):
        super().__init__()
        self.features = nn.Sequential(
            # Вход: (B, 1, 48, 48)
            nn.Conv2d(1, 16, kernel_size=3, padding=1),
            nn.ReLU(),
            nn.MaxPool2d(2, 2),

            nn.Conv2d(16, 32, kernel_size=3, padding=1),
            nn.ReLU(),
            nn.MaxPool2d(2, 2),



            nn.Conv2d(32, 64, kernel_size=3, padding=1),
            nn.ReLU(),
            nn.MaxPool2d(2, 2),
        )

        self.classifier = nn.Sequential(
            nn.Flatten(),
            nn.Linear(64 * 6 * 6, 128),
            nn.ReLU(),
            nn.Linear(128, num_classes)
        )
    ...
)
```




TRAIN LOSS

1.0686



VAL LOSS

1.1850



VAL ACCURACY

56.62%



Model №4: Improved CNN



ARCHITECTURE

```
# Block 1: (B, 1, 48, 48) -> (B, 32, 24, 24)
nn.Conv2d(1, 32, kernel_size=3, padding=1),
nn.BatchNorm2d(32),
nn.ReLU(),
nn.Conv2d(32, 32, kernel_size=3, padding=1),
nn.BatchNorm2d(32),
nn.ReLU(),
nn.MaxPool2d(2, 2),           # 48 -> 24

# Block 2: (B, 32, 24, 24) -> (B, 64, 12, 12)
nn.Conv2d(32, 64, kernel_size=3, padding=1),
nn.BatchNorm2d(64),
nn.ReLU(),
nn.Conv2d(64, 64, kernel_size=3, padding=1),
nn.BatchNorm2d(64),
nn.ReLU(),
nn.MaxPool2d(2, 2),           # 24 -> 12


# Block 3: (B, 64, 12, 12) -> (B, 128, 6, 6)
nn.Conv2d(64, 128, kernel_size=3, padding=1),
nn.BatchNorm2d(128),
nn.ReLU(),
nn.MaxPool2d(2, 2),           # 12 -> 6

self.classifier = nn.Sequential(
    nn.Flatten(),                  # (B, 128*6*6)
    nn.Dropout(0.5),
    nn.Linear(128 * 6 * 6, 256),
    nn.ReLU(),
    nn.Dropout(0.5),
    nn.Linear(256, num_classes)
)
```



TRAIN LOSS

1.2832




TRAIN ACCURACY

49.82%



VAL LOSS

1.1493



VAL ACCURACY

56.95%



Model №5: GAP CNN

ARCHITECTURE



```
# Block 1: 1x48x48 -> 32x24x24
nn.Conv2d(1, 32, kernel_size=3, padding=1),
nn.BatchNorm2d(32),
nn.ReLU(),
nn.MaxPool2d(2, 2),

# Block 2: 32x24x24 -> 64x12x12
nn.Conv2d(32, 64, kernel_size=3, padding=1),
nn.BatchNorm2d(64),
nn.ReLU(),
nn.MaxPool2d(2, 2),

# Block 3: 64x12x12 -> 128x6x6
nn.Conv2d(64, 128, kernel_size=3, padding=1),
nn.BatchNorm2d(128),
nn.ReLU(),
nn.MaxPool2d(2, 2),


# Block 4: 128x6x6 -> 256x6x6
nn.Conv2d(128, 256, kernel_size=3, padding=1),
nn.BatchNorm2d(256),
nn.ReLU(),

self.gap = nn.AdaptiveAvgPool2d(1) # -> (B, 256, 1, 1)
self.classifier = nn.Linear(256, num_classes)
```



TRAIN LOSS

0.9249




TRAIN ACCURACY

65.43%



VAL LOSS

1.1051



VAL ACCURACY



59.38%



Model №6: MiniResNet


ARCHITECTURE

- Input: **1×48×48**
- **Stem:** 3×3 Conv + BN + ReLU → 32 channels
- **3 residual stages (BasicBlock):**
- Stage 1: 32 ch, 48×48 → 24×24
- Stage 2: 64 ch, 24×24 → 12×12
- Stage 3: 128 ch, 12×12 → 6×6
- **Skip connections** with downsampling
- **Global Average Pooling**
- **Fully connected layer -> 7 classes**



TRAIN LOSS

0.6823




TRAIN ACCURACY

74.47%



VAL LOSS

1.0082



VAL ACCURACY



64.14%



Model №7: ResNet-18 (Transfer Learning)


ARCHITECTURE

- **FERResNet18 (Transfer Learning with ResNet-18)**
- Backbone: **ResNet-18** (ImageNet-pretrained, optional)
- Head: replace **final FC -> 7-class classifier**
- Input pipeline (FER 48×48 grayscale):
 • **repeat 1 -> 3 channels**
 • **resize 48×48 -> 224×224** (bilinear)
 • **ImageNet normalization** (mean/std)
- Training strategy:
- **Stage 1:** freeze backbone, train **classifier head only** (Adam, lr=1e-3)
- **Stage 2:** unfreeze all layers, **fine-tune full network** (Adam, lr=1e-4)



TRAIN LOSS

0.1976




TRAIN ACCURACY

93.55%



VAL LOSS

0.9433




VAL ACCURACY

68.21%




Evaluation & Comparison

Fine-tuned ResNet18 achieves the **best validation accuracy**, clearly outperforming simpler CNN and MLP-based architectures, highlighting the advantage of transfer learning for FER tasks




Strong performance for **Happy & Neutral** classes, while notable confusion **between similar emotions**, indicating **class overlap** in facial expressions



Evaluation & Comparison

Fine-tuned ResNet18 achieves the **best validation accuracy**, clearly outperforming simpler CNN and MLP-based architectures, highlighting the advantage of transfer learning for FER tasks



Fine-tuned ResNet18 achieves the solid test accuracy however, there is **still room for improvement**



FER TEST LOSS
1.1669

FER TEST ACC.
69.71 %

- Good convergence
- Satisfactory results
- But not perfect



Best Model Generalization Check: RAF-DB Dataset



RAF TRAIN LOSS

0.0004



RAF TRAIN ACC.

100.0%



RAF VAL LOSS

0.7009



RAF VAL ACC.

82.19%



RAF TEST LOSS

0.6347

RAF TEST ACC.

82.21%

FER TEST LOSS


2.4567

FER TEST ACC.

54.04%

The performance gap reveals **poor generalization capability. Why?**

Best Model Generalization Check: RAF-DB Dataset



! The performance gap reveals **poor generalization capability. Why?**

SOURCES OF THE ISSUES

Domain Shift

Overfitting

Class Imbalance

Catastrophic Forgetting

Conflicting Optimal Solutions

RECOMMENDATIONS

Joint Training


Freeze Lower Layers During Fine-Tuning

Hyperparameter Tuning, Regularization, and Early Stopping

Improve Data Alignment

Key Conclusions

1. CNNs significantly outperform MLPs for facial expression recognition
2. Architectural improvements and residual learning provide consistent gains
3. Pretraining is the most impactful factor for overall performance
4. Class imbalance and dataset bias remain major error sources
5. Cross-dataset generalization is limited and requires fine-tuning





Human Emotions Recognition

[GITHUB.COM/BOBLAROS](https://github.com/boblaros)